

Test Neural Network

October 29, 2018

```
In [1]: from mnist import MNIST
import numpy as np
from matplotlib import pyplot as plt
import types
from sklearn.utils import shuffle
from sklearn.metrics import confusion_matrix
import pandas as pd

In [2]: def load_data():
    mnist_data = MNIST('./')
    train_img, train_label = mnist_data.load_training()
    test_img, test_label = mnist_data.load_testing()
    return np.array(train_img), np.array(train_label, dtype="uint8"), np.array(test_img)

    # Execute function
    train_img, train_label, test_img, test_label = load_data()

In [3]: def vectorize_label(label, num_class):
    return np.eye(num_class)[label.reshape(-1)]

    # Execute function
    NUM_CLASS = 10
    train_label = vectorize_label(train_label, NUM_CLASS)
    test_label = vectorize_label(test_label, NUM_CLASS)

In [4]: def initialize_weight(width, height):
    range_ = 0.12
    return np.random.rand(width,height) * range_ * 2 - range_

In [5]: def derivative_linear(X):
    return np.ones_like(X)

In [6]: def relu(z):
    return z * (z > 0) * 1

In [7]: def derivative_relu(X):
    return (X>0) * 1

In [8]: def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

```

In [9]: def derivative_sigmoid(X):
        return X * (1 - X)

In [10]: def softmax(z):
        activate = np.exp(z - z.max(axis=0))
        return activate / activate.sum(axis=0)

In [11]: def cost_softmax(activation, ground_truth, all_W=0, lambdaa=0):
        m = np.size(activation, 1) # Number of example
        regularized_term_cost = 0
        if isinstance(all_W, list):
            for single_W in all_W:
                W_sliced = single_W.copy()
                W_sliced[:,0] = 0 # Zeroes out bias weight
                regularized_term_cost += lambdaa / 2 / m * (W_sliced ** 2).sum();
        cost = (-ground_truth * np.log(activation)).sum() / m + regularized_term_cost;
        return cost

In [12]: def grad_softmax(activation, ground_truth, W, X, derivative_method=None, lambdaa = 0):
        m = np.size(activation, 1) # Number of example
        W_sliced = W.copy()
        if W_sliced is not 0:
            W_sliced[:,0] = 0 # Zeroes out bias weight

        err = (activation - ground_truth) / m
        regularized_term_grad = lambdaa / m * W_sliced
        grad = err.dot(X.T) + regularized_term_grad
        if isinstance(derivative_method, types.FunctionType):
            err = W.T.dot(err) * derivative_method(X)
        return grad, err

In [13]: def softmax_classifier(all_W, X, Y, derivative_method=None, lambdaa = 0):
        if isinstance(all_W, list):
            W = all_W[len(all_W) - 1] # Softmax weight

            # Forward pass
            activation = softmax(W.dot(X));
            pred = activation
            ground_truth = Y

            # Compute cost
            cost = cost_softmax(activation, ground_truth, all_W, lambdaa)

            # Compute gradient and error
            grad, err = grad_softmax(activation, ground_truth, W, X, derivative_method, 1)
            return cost, grad, err, pred

In [14]: def numerical_softmax_check(W, X, Y, lambdaa=0):
        eps = 1e-6

```

```

grad = np.zeros_like(W)
for i in range(np.size(W,0)):
    for j in range(np.size(W,1)):
        temp = W[i,j]
        W[i,j] = temp + eps
        cost1, grad1, err1, pred1 = softmax_classifier([W],X,Y,lambdadaa=lambdadaa)
        W[i,j] = temp - eps
        cost2, grad2, err2, pred2 = softmax_classifier([W],X,Y,lambdadaa=lambdadaa)
        W[i,j] = temp
        grad[i,j] = (cost1 - cost2) / 2 / eps
return grad

```

```

In [15]: def backprop(err, W, X, derivative_method=None, lambdadaa = 0):
    m = np.size(X, 1) # Number of example
    W_sliced = W.copy()
    if W_sliced is not 0:
        W_sliced[:,0] = 0 # Zeroes out bias weight
    regularized_term = lambdadaa / m * W_sliced
    err = err[1:,:] # Shed virtual err
    grad = err.dot(X.T) + regularized_term
    if isinstance(derivative_method, types.FunctionType):
        err = W.T.dot(err) * derivative_method(X)
    return grad, err

```

```

In [16]: W = initialize_weight(2,3)
X = initialize_weight(3,3)
ground_truth = np.zeros((2,3))
ground_truth[0,:] = 1
# Check softmax grad function without regularized term
grad = numerical_softmax_check(W,X,ground_truth)
cost2, grad2, err2, pred2 = softmax_classifier([W],X,ground_truth)
# Check softmax grad function with regularized term = 1
grad3 = numerical_softmax_check(W,X,ground_truth, lambdadaa=1)
cost4, grad4, err4, pred4 = softmax_classifier([W],X,ground_truth, lambdadaa=1)
print("Relative difference without regularization\n",np.linalg.norm(grad - grad2) / np
print("Relative difference with regularization = 1\n",np.linalg.norm(grad3 - grad4) /

```

Relative difference without regularization

1.2726672348055012e-09

Relative difference with regularization = 1

7.893984534802521e-10

```

In [17]: def neural_network(all_W, train_data, train_label, lambdadaa=0):
    m = np.size(train_data, 1) # Number of training example
    list_activation = []
    list_grad = []

    row_of_one = np.ones((1,m))

```

```

activation = train_data
activation = np.vstack((row_of_one, activation))
list_activation.append(activation)

# Forward pass
for i in range(len(all_W) - 1):
    Z = all_W[i].dot(activation)
    activation = sigmoid(Z)
    activation = np.vstack((row_of_one, activation))
    list_activation.append(activation)
cost, grad, err, pred = softmax_classifier(all_W, activation, train_label, derivative_sigmoid, lambdaa)
list_activation.append(pred)
list_grad.append(grad)

# Backward pass
for i in range(len(all_W) - 2, -1, -1):
    grad, err = backprop(err, all_W[i], list_activation[i], derivative_sigmoid, lambdaa)
    list_grad.append(grad)

return list_grad, list_activation, cost;

```

```

In [18]: def numerical_grad_check(all_W, X, Y, lambdaa=0):
    eps = 1e-6
    list_grad = []
    for k in range(len(all_W)-1,-1,-1):
        W = all_W[k]
        grad = np.zeros_like(W)
        for i in range(np.size(W,0)):
            for j in range(np.size(W,1)):
                temp = W[i,j]
                W[i,j] = temp + eps
                list_grad, list_activation, cost = neural_network(all_W, data, label, lambdaa)
                W[i,j] = temp - eps
                list_grad2, list_activation2, cost2 = neural_network(all_W, data, label, lambdaa)
                grad[i,j] = (cost - cost2) / 2 / eps
                W[i,j] = temp
            list_grad.append(grad)
    return list_grad

```

```

In [19]: # Check gradient backprop
data = np.random.rand(784, 10)
label = np.zeros((10, 10))
label[0,:] = 1
W1 = initialize_weight(25, 785)
W2 = initialize_weight(15,26)
W3 = initialize_weight(10,16)
all_W = [W1, W2, W3]
lambdaa = 1

```

```

%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
list_grad, list_activation, cost = neural_network(all_W, data, label, lambdab)
list_grad2 = numerical_grad_check(all_W, data, label, lambdab)

print("Relative difference first layer\n", np.linalg.norm(list_grad[2] - list_grad2[2])
print("Relative difference second layer\n", np.linalg.norm(list_grad[1] - list_grad2[1])
print("Relative difference third layer\n", np.linalg.norm(list_grad[0] - list_grad2[0])

```

```

Relative difference first layer
5.1520981516639004e-08
Relative difference second layer
1.9481216244011975e-08
Relative difference third layer
2.3621286933308203e-09

```

```

In [20]: def training(all_W, train_data, train_label, batch_size=30, learning_rate=0.001, num_loops=1000):
    m = np.size(train_data, 1)
    average_cost = 0
    times = 0
    for i in range(num_loops):
        # Shuffle data randomly
        if shuffle is not 0:
            print("Shuffle")
            permutation = np.random.permutation(m)
            train_data = train_data[:,permutation]
            train_label = train_label[:, permutation]
            print("Done")

        for j in range(0,m,batch_size):
            list_gradient, list_activation, cost = neural_network(all_W, train_data[j:j+batch_size,:])
            list_gradient.reverse()
            for j in range(len(all_W)): # Update all W
                all_W[j] -= learning_rate * list_gradient[j]
            average_cost += cost
            times += 1
        print("Iterate ", i+1, " loss: ", average_cost / times)
        average_cost = 0
        times = 0

```


```



In [21]: train_input = (train_img/255).T # Train data shape: 784 x 60000
train_truth = train_label.T # Train label shape: 10 x 60000
# First layer logistic - shape: 25 x 785
W1 = initialize_weight(25, 785)
# Second layer logistic - shape: 15 x 26
W2 = initialize_weight(15, 26)
# Final layer - softmax classifier - shape: 10 x 16
W3 = initialize_weight(10, 16)
all_W = [W1, W2, W3]



```



```
In [22]: first_layer = all_W[0][:,1:]
print(first_layer.shape)
i = 0
for row in first_layer:
    img = np.reshape(row,(28,28))
    plt.imshow(img)
    plt.show()
```



(25, 784)











































```
In [23]: training(all_W, train_input, train_truth, batch_size=30, num_loop=3, learning_rate=0.)
```


```
Iterate 1 loss: 2.067492991886211
```



```
Iterate 2 loss: 0.7916308968892606
```



Iterate 3 loss: 0.4392944543728588



```
In [24]: first_layer = all_W[0][:,1:]
print(first_layer.shape)
i = 0
for row in first_layer:
    img = np.reshape(row,(28,28))
    plt.imshow(img)
    plt.show()
```



(25, 784)











































```
In [25]: def plot_confusion_matrix(df_confusion, title='Confusion matrix', cmap=plt.cm.gray_r)
plt.matshow(df_confusion, cmap=cmap) # imshow
plt.title(title)
```

```

plt.colorbar()
tick_marks = np.arange(len(df_confusion.columns))
plt.xticks(tick_marks, df_confusion.columns, rotation=45)
plt.yticks(tick_marks, df_confusion.index)
#plt.tight_layout()
plt.ylabel(df_confusion.index.name)
plt.xlabel(df_confusion.columns.name)

```

```

In [26]: test_img_input = (test_img/255).T
test_label_truth = test_label.T

```

```

list_grad, list_activation, cost = neural_network(all_W, test_img_input, test_label_t
pred = list_activation[len(all_W)].argmax(0)
answer = test_label_truth.argmax(0)
compare_vector = (pred == answer)
num_correct = np.sum(compare_vector)
m = len(compare_vector)
print("Accuracy: ", num_correct / m * 100,"%")
print("Cost: ", cost)
print()
print("Confusion matrix: ")
y_actu = pd.Series(answer, name='Actual')
y_pred = pd.Series(pred, name='Predicted')
df_confusion = pd.crosstab(y_actu, y_pred)
print(df_confusion)
print()
print("Normalized confusion matrix: ")
df_conf_norm = round(df_confusion / df_confusion.sum(axis=1) * 100, 1)
print(df_conf_norm)

```

```

Accuracy: 89.57000000000001 %
Cost: 0.3746792211602384

```

Confusion matrix:

Predicted \ Actual	0	1	2	3	4	5	6	7	8	9
0	944	0	3	1	4	17	8	1	1	1
1	0	1097	5	5	0	1	3	1	22	1
2	14	17	888	24	14	5	30	7	29	4
3	5	8	19	882	1	51	1	21	14	8
4	2	0	2	0	906	0	16	1	5	50
5	16	2	8	43	12	754	21	4	26	6
6	28	2	6	0	14	15	890	0	3	0
7	2	26	13	12	6	2	0	874	2	91
8	4	11	5	21	21	42	15	4	827	24
9	4	2	1	7	50	21	0	23	6	895

Normalized confusion matrix:

Predicted Actual	0	1	2	3	4	5	6	7	8	9
0	96.3	0.0	0.3	0.1	0.4	1.9	0.8	0.1	0.1	0.1
1	0.0	96.7	0.5	0.5	0.0	0.1	0.3	0.1	2.3	0.1
2	1.4	1.5	86.0	2.4	1.4	0.6	3.1	0.7	3.0	0.4
3	0.5	0.7	1.8	87.3	0.1	5.7	0.1	2.0	1.4	0.8
4	0.2	0.0	0.2	0.0	92.3	0.0	1.7	0.1	0.5	5.0
5	1.6	0.2	0.8	4.3	1.2	84.5	2.2	0.4	2.7	0.6
6	2.9	0.2	0.6	0.0	1.4	1.7	92.9	0.0	0.3	0.0
7	0.2	2.3	1.3	1.2	0.6	0.2	0.0	85.0	0.2	9.0
8	0.4	1.0	0.5	2.1	2.1	4.7	1.6	0.4	84.9	2.4
9	0.4	0.2	0.1	0.7	5.1	2.4	0.0	2.2	0.6	88.7